

Confluent Cloud Clients Python Library

The Confluent Cloud Clients Python Library provides a set of clients for interacting with Confluent Cloud REST APIs. The library includes clients for:

- **Flink**
- **Kafka**
- **Schema Registry**
- **Tableflow**
- **Metrics**

Note: *This library is in active development and is subject to change. It covers only the methods I have needed so far. If you need a method that is not covered, please feel free to open an issue or submit a pull request.*

Table of Contents

- **1.0 Library Clients**
 - **1.1 Flink Client**
 - **1.2 Kafka Client**
 - **1.3 Schema Registry Client**
 - **1.4 Tableflow Client**
 - **1.5 Metrics Client**
- **2.0 Unit Tests**
 - **2.1 Flink Client**
 - **2.2 Kafka Client**
 - **2.3 Schema Registry Client**
 - **2.4 Tableflow Client**
 - **2.5 Metrics Client**
- **3.0 Installation**
- **4.0 Resources**
 - **4.1 Architecture Design Records (ADRs)**
 - **4.2 API Documentation**
 - **4.3 Flink Resources**
 - **4.4 Tableflow Resources**
 - **4.5 Metrics Resources**
 - **4.6 Other Resources**

1.0 Library Clients

1.1 Flink Client

The **Flink Client** provides the following methods:

- `delete_statement`
- `delete_statements_by_phase`
- `drop_table`

Note: *"The `drop_table` method will drop the table and all associated statements, including the backing Kafka Topic and Schemas."*

- `get_compute_pool`
- `get_compute_pool_list`
- `get_statement_list`
- `stop_statement`

Note: "Confluent Cloud for Apache Flink enforces a **30-day** retention for statements in terminal states."

- `submit_statement`
- `update_statement`
- `update_all_sink_statements`

1.2 Kafka Client

The **Kafka Client** provides the following methods:

- `delete_kafka_topic`
- `kafka_topic_exist`

1.3 Schema Registry Client

The **Schema Registry Client** provides the following methods:

- `convert_avro_schema_into_string`
- `delete_kafka_topic_key_schema_subject`
- `delete_kafka_topic_value_schema_subject`
- `get_global_topic_subject_compatibility_level`
- `get_topic_subject_compatibility_level`
- `get_topic_subject_latest_schema`
- `register_topic_subject_schema`
- `set_topic_subject_compatibility_level`

1.4 Tableflow Client

The **Tableflow Client** provides the following methods:

- `get_tableflow_topic`
- `get_tableflow_topic_table_path`

1.5 Metrics Client

The **Metrics Client** provides the following methods:

- `get_topic_total`

Note: "The `get_topic_total` method can be used to get the total bytes or total records for a Kafka Topic. It requires an additional parameter to specify the metric type."

◦ Metric Types:

- `RECEIVED_BYTES`
- `RECEIVED_RECORDS`

- `get_topic_daily_aggregated_totals`

Note: "The `get_topic_daily_aggregated_totals` method can be used to get the daily aggregated totals for a Kafka Topic within a rolling window of the last 7 days. It requires an additional parameter to specify the metric type."

◦ Metric Types:

- `RECEIVED_BYTES`
- `RECEIVED_RECORDS`

2.0 Unit Tests

The library includes unit tests for each client. The tests are located in the `tests` directory. To use them, you must clone the repo locally:

```
git clone https://github.com/j3-signalroom/cc-clients-python_lib.git
```

Since this project was built using **uv**, please install it, and then run the following command to install all the project dependencies:

```
uv sync
```

Then within the **tests** directory, create the **.env** file and add the following environment variables, filling them with your Confluent Cloud credentials and other required values:

```
BOOTSTRAP_SERVER_CLOUD_PROVIDER=
BOOTSTRAP_SERVER_CLOUD_REGION=
BOOTSTRAP_SERVER_ID=
CLOUD_PROVIDER=
CLOUD_REGION=
COMPUTE_POOL_ID=
CONFLUENT_CLOUD_API_KEY=
CONFLUENT_CLOUD_API_SECRET=
ENVIRONMENT_ID=
FLINK_API_KEY=
FLINK_API_SECRET=
FLINK_CATALOG_NAME=
FLINK_DATABASE_NAME=
FLINK_STATEMENT_NAME=
FLINK_TABLE_NAME=
FLINK_URL=
KAFKA_API_KEY=
KAFKA_API_SECRET=
KAFKA_CLUSTER_ID=
KAFKA_TOPIC_NAME=
ORGANIZATION_ID=
PRINCIPAL_ID=
QUERY_START_TIME=
QUERY_END_TIME=
SCHEMA_REGISTRY_API_KEY=
SCHEMA_REGISTRY_API_SECRET=
SCHEMA_REGISTRY_URL=
TABLEFLOW_API_KEY=
TABLEFLOW_API_SECRET=
```

Note: The **QUERY_START_TIME** and **QUERY_END_TIME** environment variables should be in the format **YYYY-MM-DDTHH:MM:SS**, for example, **2025-09-01T00:00:00**.

2.1 Flink Client

To run a specific test, use one of the following commands:

Unit Test	Command
Delete a Flink Statement	uv run pytest -s tests/test_flink_client.py::test_delete_statement
Delete all Flink Statements by Phase	uv run pytest -s tests/test_flink_client.py::test_delete_statements_by_phase

Unit Test	Command
Get list of the all the Statements	<code>uv run pytest -s tests/test_flink_client.py::test_get_statement_list</code>
Submit a Flink Statement	<code>uv run pytest -s tests/test_flink_client.py::test_submit_statement</code>
Get Compute Pool List	<code>uv run pytest -s tests/test_flink_client.py::test_get_compute_pool_list</code>
Get Compute Pool	<code>uv run pytest -s tests/test_flink_client.py::test_get_compute_pool</code>
Stop a Flink Statement	<code>uv run pytest -s tests/test_flink_client.py::test_stop_statement</code>
Update a Flink Statement	<code>uv run pytest -s tests/test_flink_client.py::test_update_statement</code>
Update all the Sink Statements	<code>uv run pytest -s tests/test_flink_client.py::test_update_all_sink_statements</code>
Drop a Flink Table along with any associated statements, including the backing Kafka Topic and Schemas	<code>uv run pytest -s tests/test_flink_client.py::test_drop_table</code>

Otherwise, to run all the tests, use the following command:

```
uv run pytest -s tests/test_flink_client.py
```

Note: The tests are designed to be run in a specific order. If you run them out of order, you may encounter errors. The tests are also designed to be run against a Confluent Cloud environment. If you run them against a local environment, you may encounter errors.

2.2 Kafka Client

To run a specific test, use one of the following commands:

Unit Test	Command
Delete a Kafka Topic	<code>uv run pytest -s tests/test_kafka_client.py::test_delete_kafka_topic</code>
Checks if a Kafka Topic Exist	<code>uv run pytest -s tests/test_kafkaclient.py::test_kafka_topic_exist</code>

Otherwise, to run all the tests, use the following command:

```
uv run pytest -s tests/test_kafka_client.py
```

Note: The tests are designed to be run in a specific order. If you run them out of order, you may encounter errors. The tests are also designed to be run against a Confluent Cloud environment. If you run them against a local environment, you may encounter errors.

2.3 Schema Registry Client

To run a specific test, use one of the following commands:

Unit Test	Command
Get the Subject Compatibility Level	<code>uv run pytest -s tests/test_schema_registry_client.py::test_get_subject_compatibility_level</code>
Delete the Kafka Topic Key Schema Subject	<code>uv run pytest -s tests/test_schema_registry_client.py::test_delete_kafka_topic_key_schema_subject</code>
Delete the Kafka Topic Value Schema Subject	<code>uv run pytest -s tests/test_schema_registry_client.py::test_delete_kafka_topic_value_schema_subject</code>

Otherwise, to run all the tests, use the following command:

```
uv run pytest -s tests/test_schema_registry_client.py
```

Note: The tests are designed to be run in a specific order. If you run them out of order, you may encounter errors. The tests are also designed to be run against a Confluent Cloud environment. If you run them against a local environment, you may encounter errors.

2.4 Tableflow Client

To run a specific test, use one of the following commands:

Unit Test	Command
Get the Tableflow Topic	<code>uv run pytest -s tests/test_tableflow_client.py::test_get_tableflow_topic</code>
Get the Tableflow Topic Table Path	<code>uv run pytest -s tests/test_tableflow_client.py::test_get_tableflow_topic_table_path</code>

Otherwise, to run all the tests, use the following command:

```
uv run pytest -s tests/test_tableflow_client.py
```

Note: The tests are designed to be run in a specific order. If you run them out of order, you may encounter errors. The tests are also designed to be run against a Confluent Cloud environment. If you run them against a local environment, you may encounter errors.

2.5 Metrics Client

To run a specific test, use one of the following commands:

Unit Test	Command
Get the Topic Total Bytes	<code>uv run pytest -s tests/test_metrics_client.py::test_get_topic_total_bytes</code>

Unit Test	Command
Get the Topic Total Records	<code>uv run pytest -s tests/test_metrics_client.py::test_get_topic_total_records</code>
Get the Topic Daily Aggregated Totals Bytes	<code>uv run pytest -s tests/test_metrics_client.py::test_get_topic_daily_aggregated_totals_bytes</code>
Get the Topic Daily Aggregated Totals Records	<code>uv run pytest -s tests/test_metrics_client.py::test_get_topic_daily_aggregated_totals_records</code>

Otherwise, to run all the tests, use the following command:

```
uv run pytest -s tests/test_metrics_client.py
```

Note: *The tests are designed to be run in a specific order. If you run them out of order, you may encounter errors. The tests are also designed to be run against a Confluent Cloud environment. If you run them against a local environment, you may encounter errors.*

3.0 Installation

Install the Confluent Cloud Clients Python Library using **pip**:

```
pip install cc-clients-python-lib
```

Or, using **uv**:

```
uv add cc-clients-python-lib
```

4.0 Resources

4.1 Architecture Design Records (ADRs)

- [001 Architectural Design Record \(ADR\): Drop Table Plus](#)

4.2 API Documentation

- [Flink SQL REST API for Confluent Cloud for Apache Flink](#)
- [Kafka REST APIs for Confluent Cloud](#)
- [Confluent Cloud APIs - Topic \(v3\)](#)
- [Confluent Cloud Schema Registry REST API Usage](#)

4.3 Flink Resources

- [CCAF State management](#)
- [Monitor and Manage Flink SQL Statements in Confluent Cloud for Apache Flink](#)
- [DROP TABLE Statement in Confluent Cloud for Apache Flink](#)

4.4 Tableflow Resources

- [Tableflow Topics \(tableflow/v1\)](#)

4.4 Tableflow Resources

- [Tableflow Topics \(tableflow/v1\)](#)

4.5 Metrics Resources

- [Confluent Cloud Metrics API Version 2 Reference](#)
- [Confluent Cloud Metrics API: Metrics Reference](#)
- [Confluent Cloud Metrics](#)

4.6 Other Resources

- [How to programmatically pause and resume a Flink statement](#)
- [How to programmatically pause and resume a Flink statement REDUX](#)